Navigating China Internet: DeepSeek V4 launch — implications for China AI models/cloud/data centers
领航中国互联网：DeepSeek V4 发布——对中国 AI 模型/云/数据中心的影响

24 April 2026 | 9:01PM HKT | Research | Equity| By Ronald Keung, CFA and others
2026 年 4 月 24 日 | 晚上 9:01 HKT | 研究 | 股票 | 作者：Ronald Keung, CFA 等

PDF Share Bookmark More

Compared with the DeepSeek defining moment last year, we view DeepSeek's latest V4 model launch as a continuation of its focus in computing efficiency and open-source approach and highlight DeepSeek V4's architectural upgrades (vs. V3) that enable: **1) significantly less memory needed for long context window,** where V4 supports 1M context window that is comparable to US state-of-the-art (SOTA) models but at a fraction of memory/KV cache needed (at 7-10% of V3 model, for its 1.6T Pro and 284bn parameter Flash models), thus could facilitate further proliferation of agentic applications given highly affordable costs, and **2) adoption of domestic chips**, where DeepSeek expects significant price cuts for its Pro models from 2H 2026, following its anticipated mass availability of Huawei's Ascend 950 super nodes. **Implications for China AI models/cloud/data centers:** We note the rise in competition amongst Chinese AI models, given the recent accelerated pace of new model launches against DeepSeek's open-source V4 (Kimi K2.6, Alibaba's Qwen3.6-Max, Tencent's Hy3 preview, Xiaomi V2.5, and potential for MiniMax's M3/Hailuo launches in May, GSe), where coding/task completion success rates and multi-modal will likely be key differentiating factors for model pricing power ahead. We continue to rank **Cloud & Data Centers** as our #1 preferred sub-sector (with key ideas: GDS, VNET, Alibaba and Kingsoft Cloud, see report) on proliferation of AI token demand with improving cloud pricing driven by enterprise/AI agent growth and consumer AI assistants. We believe improvement in computing cost efficiencies could drive further room for wider adoption/exploration/proliferation of AI applications.

与去年 DeepSeek 的高光时刻相比，我们认为 DeepSeek 最新发布的 V4 模型延续了其对计算效率和开源路径的专注。我们重点关注 DeepSeek V4 相比 V3 的架构升级，这些升级实现了：1) 长上下文窗口所需的内存显著减少，V4 支持 100 万（1M）上下文窗口，与美国最先进（SOTA）模型相当，但所需的内存/KV 缓存仅为极小部分（其 1.6T Pro 模型和 284bn 参数 Flash 模型仅需 V3 模型的 7-10%），鉴于极具性价比的成本，这可能促进智能体（agentic）应用的进一步普及；2) 采用国产芯片，DeepSeek 预计在华为昇腾 950 超级节点大规模可用后，其 Pro 模型将从 2026 年下半年开始大幅降价。对中国 AI 模型/云/数据中心的影响：我们注意到中国 AI 模型之间的竞争日益加剧，近期新模型发布节奏加快，以应对 DeepSeek 的开源 V4（包括 Kimi K2.6、阿里巴巴的通义千问 Qwen3.6-Max、腾讯混元 Hy3 预览版、小米 V2.5，以及 MiniMax 可能在 5 月发布的 M3/海螺，高盛预测），其中代码编写/任务完成成功率和多模态能力可能是未来模型定价权的关键差异化因素。我们继续将云与数据中心列为首选子行业（核心标的：万国数据、世纪互联、阿里巴巴和金山云，详见报告），原因在于 AI Token 需求激增，且在企业/AI 智能体增长和消费者 AI 助手的推动下，云定价环境有所改善。我们认为计算成本效率的提升将为 AI 应用的更广泛采用、探索和普及提供进一步空间。

Key highlights of DeepSeek V4:

DeepSeek V4 的核心亮点：

- **DeepSeek released its latest open-sourcing DeepSeek-V4 Preview on Apr 24, 2026,** in two versions (Pro and Flash), where the Pro version is a flagship-scale model featuring 1.6T parameters (with 49B activated), while the Flash version is relatively small with 284 billion parameters (and 13B activated).

  DeepSeek 于 2026 年 4 月 24 日发布了其最新的开源预览版 DeepSeek-V4，包含两个版本（Pro 和 Flash）。其中 Pro 版本是旗舰级模型，拥有 1.6 万亿参数（激活参数为 490 亿）；而 Flash 版本规模相对较小，拥有 2840 亿参数（激活参数为 130 亿）。

  We note DeepSeek continues to be foundation text focused, while internet mega-caps and other independent model players like MiniMax are more multi/full-modal focus in their model approach, in their ongoing pursuit of AGI.

  我们注意到，DeepSeek 继续专注于基础文本模型，而互联网巨头以及 MiniMax 等其他独立模型厂商在追求通用人工智能（AGI）的过程中，其模型路径更倾向于多模态或全模态。

- **1M context window for both Pro and Flash models, via new architectural breakthroughs:** Both Pro and Flash models support an ultra-long context window of 1M tokens, with substantially lower inference FLOPS and KV-cache footprint vs. DeepSeek-V3.2, helped by 3 key upgrades. Besides DeepSeek MoE and Multi-Token Prediction, DeepSeek V4 incorporates 3 key upgrades in architecture and optimizations:

  通过全新的架构突破，Pro 和 Flash 模型均支持 100 万（1M）上下文窗口：得益于三项关键升级，Pro 和 Flash 模型均支持 100 万 token 的超长上下文窗口，且与 DeepSeek-V3.2 相比，推理算力（FLOPS）需求和 KV 缓存占用大幅降低。除了 DeepSeek MoE 和多 token 预测（Multi-Token Prediction）之外，DeepSeek V4 在架构和优化方面还引入了三项核心升级：

  - **Hybrid attention architecture combines CSA and HCA.** Compressed Sparse Attention (CSA) first compresses the KV cache along the sequence dimension and then applies sparse attention, while Heavily Compressed Attention (HCA) applies more aggressive KV compression but keeps dense attention. As a result, V4 reduces the amount of temporary memory the model needs to keep for long inputs.

    混合注意力架构结合了 CSA 和 HCA。压缩稀疏注意力（CSA）首先沿序列维度压缩 KV 缓存，然后应用稀疏注意力；而重度压缩注意力（HCA）则采用更激进的 KV 压缩方式，但保持稠密注意力。因此，V4 减少了模型在处理长输入时需要保留的临时内存量。

  - **mHC improves training reliability:** DeepSeek-V4 adds mHC to make the model more stable as information passes through many layers.

    mHC 提升了训练可靠性：DeepSeek-V4 加入了 mHC，使得模型在信息经过多层传递时更加稳定。

  - **Muon optimizer:** DeepSeek-V4 uses Muon as its main training optimizer, while still using AdamW for some smaller/special modules. In simple terms, Muon is the training method that helps the model learn in a more stable way as V4 has a more complex architecture than V3.

    Muon 优化器：DeepSeek-V4 使用 Muon 作为主要训练优化器，同时在一些较小或特殊的模块中仍沿用 AdamW。简单来说，由于 V4 的架构比 V3 更复杂，Muon 这种训练方法能帮助模型以更稳定的方式进行学习。

  - These updates result in **meaningful efficiency gains versus DeepSeek-V3.2.** At 1M context, V4-Pro reportedly uses **27% of DeepSeek V3.2's single-token inference FLOPs** and **10% of its KV cache**; V4-Flash is reported at **10% FLOPs** and **7% KV cache**. This suggests efficiency gains for long-context workloads, where the likely use-case relevance is long-horizon tasks rather than ordinary short prompts.

    这些更新带来了相较于 DeepSeek-V3.2 显著的效率提升。据报道，在 1M 上下文长度下，V4-Pro 的单 token 推理 FLOPs 仅为 DeepSeek V3.2 的 27%，KV cache 仅为 10%；V4-Flash 的 FLOPs 更是低至 10%，KV cache 仅为 7%。这表明长上下文工作负载的效率得到了大幅提升，其应用场景更倾向于长程任务，而非普通的短提示词。

Implications for China AI models/cloud/data centers:

对中国 AI 模型/云服务/数据中心的影响：

- With DeepSeek-V4's competitive model capabilities, **we continue to rank Cloud & Data Centers as our preferred sub-sector (#1) (with key ideas: GDS, VNET, Alibaba and Kingsoft Cloud)** on continued proliferation of AI token demand with improving cloud/token pricing power driven by enterprise/AI agent growth and consumer AI assistants. We believe further improvement in computing cost efficiencies could drive further room for wider adoption/exploration/proliferation of AI applications.

  凭借 DeepSeek-V4 极具竞争力的模型能力，我们继续将云与数据中心列为首选子行业（排名第一）（核心标的包括：万国数据、世纪互联、阿里巴巴和金山云）。这主要基于企业级/AI 智能体增长以及消费者 AI 助手推动下，AI Token 需求持续激增，且云/Token 定价能力不断提升。我们认为，计算成本效率的进一步提高，将为 AI 应用的更广泛采用、探索和普及提供更大空间。

- For **AI models**, given the recent accelerated pace of new model launches against DeepSeek's open-source V4 (Kimi K2.6, Alibaba's Qwen3.6-Max, Tencent's Hy3 preview, Xiaomi V2.5, and potential for MiniMax's M3/Hailuo launches in May, GSe), where coding (noting Knowledge Atlas/Zhipu's GLM model is highly ranked), task completion success rates and multi-modal (ByteDance, Alibaba and MiniMax) will likely be key differentiating factors for model pricing power ahead.

  在 AI 模型方面，鉴于近期新模型发布节奏加快，正面对标 DeepSeek 的开源 V4（包括 Kimi K2.6、阿里巴巴的 Qwen3.6-Max、腾讯混元 3.6 预览版、小米 V2.5，以及 MiniMax 可能在 5 月发布的 M3/海螺模型，高盛预测），代码能力（注：智谱 GLM 模型排名靠前）、任务完成成功率以及多模态能力（字节跳动、阿里巴巴和 MiniMax）可能成为未来模型定价权的关键差异化因素。

- We believe independent players' competitive advantages vs. internet mega-caps rest on their high organization efficiency and decision-making processes in identifying next key AI model developments, e.g., MiniMax's highly efficient model design and inference enables the company to achieve 40% GPM (GSe) for its foundation text API channel, even at highly competitive API pricing levels.

  我们认为，独立厂商相对于互联网巨头的竞争优势在于其极高的组织效率和决策流程，能够快速识别下一个关键的 AI 模型发展方向。例如，MiniMax 高效的模型设计和推理能力使其基础文本 API 渠道即使在极具竞争力的 API 定价水平下，仍能实现 40% 的毛利率（高盛预测）。

- For **internet mega-caps**, we believe they are best positioned to capture the AI infrastructure/cloud opportunity given strong underlying operating cash flows from their core business, while separate standalone incentive schemes for AI chip/model teams (e.g., Doubao AI team already has separately incentive programs) will likely be needed to incentivise/retain top AI talent vs. independent AI native players.

  对于互联网巨头而言，我们认为凭借其核心业务产生的强大底层经营现金流，他们最适合捕捉 AI 基础设施/云市场的机遇。同时，为了与独立 AI 原生厂商竞争并吸引/留住顶尖 AI 人才，巨头们可能需要为 AI 芯片/模型团队设立独立的激励方案（例如，豆包 AI 团队已经拥有独立的激励计划）。

  We also notenews reports that Tencent and Alibaba are potentially in talks to invest in DeepSeek at over US$20bn valuation, compared with last close market caps for Knowledge Atlas (Zhipu)/MiniMax at US$53bn/31bn.

  我们还注意到有新闻报道称，腾讯和阿里巴巴可能正洽谈以超过 200 亿美元的估值投资 DeepSeek，而智谱 AI (Knowledge Atlas) 和 MiniMax 的最近一轮市场估值分别为 53 亿美元和 31 亿美元。

Related research: 相关研究：

Navigating China Internet: What to do from here & key focuses post results season; addressing key AI debates, April 1, 2026

中国互联网行业导航：后续操作建议及业绩季后的核心关注点；应对 AI 关键争论，2026 年 4 月 1 日

Tencent Holdings (0700.HK): Hy3 preview marks a key development in Tencent's AI revamp; Buy, April 23, 2026

腾讯控股(0700.HK)：2026 年上半年预览标志着腾讯 AI 架构重组的关键进展；买入，2026 年 4 月 23 日

Xiaomi Corp. (1810.HK): MiMo-V2.5 series release; Accelerating model update pace paves way for commercialization and expanding AI applications; Buy, April 24, 2026

小米集团(1810.HK)：MiMo-V2.5 系列发布；加速模型更新步伐，为商业化和扩大 AI 应用铺平道路；买入，2026 年 4 月 24 日

MiniMax Group (0100.HK): Global full-modal AI company reaching hypergrowth stage; initiate at Neutral on valuation, Feb 23, 2026

MiniMax 集团(0100.HK)：全球全模态 AI 公司进入高速增长阶段；基于估值给予中性评级，2026 年 2 月 23 日

The authors would like to thank Iris Xiao for her contribution to this report.

作者感谢 Iris Xiao 对本报告的贡献。

Exhibit 1: Our quarterly updated sub-sector preference within China Internet (Apr 2026 update)

图表 1：我们每季度更新的中国互联网子行业偏好（2026 年 4 月更新）

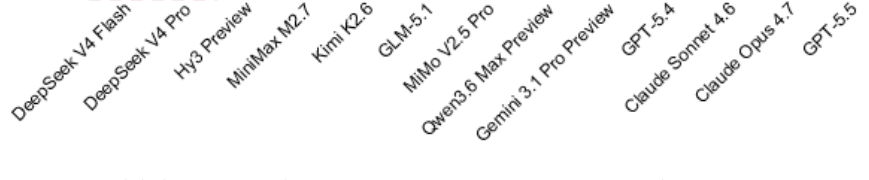

Sep 2025 update	Dec 2025 update	Jan 2026 update	Mar 2026 update
Games	Cloud & data center	Cloud & data center	Cloud & data center
Mobility	Games	Games & Entertainment	eCommerce & mobility
Cloud & data center	Mobility	AI model	Games & Entertainment
eCommerce & local service	eCommerce & local service	eCommerce & mobility	AI model

Source: Goldman Sachs Global Investment Research
来源：高盛全球投资研究部

DeepSeek-V4 pricing/performance comparison with counterparts
DeepSeek-V4 与同类产品的价格/性能对比

Exhibit 2: DeepSeek V4 continues to show pricing competitiveness with V4 Pro price expected to further decline with Huawei-backed domestic compute capacity ramps in 2H26

图表 2：DeepSeek V4 持续展现价格竞争力，随着华为支持的国内算力在 2026 年下半年逐步释放，预计 V4 Pro 的价格将进一步下降

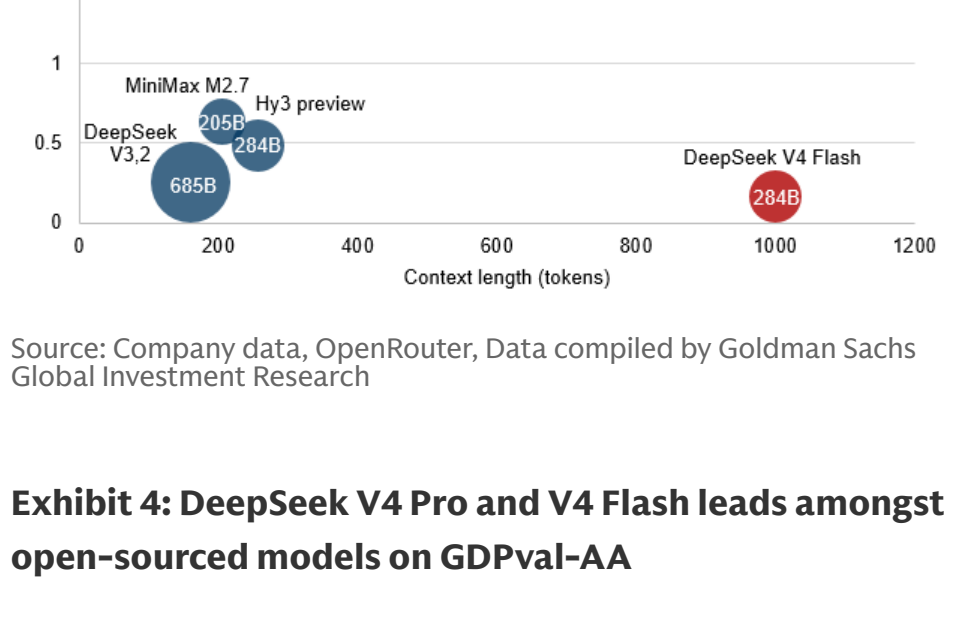

Source: Artificial Analysis, Company data, Data compiled by Goldman Sachs Global Investment Research

Exhibit 3: DeepSeek V4 pushes 1M context with better cost efficiency amongst open-source models
Mimo V2.5 Pro to be open-sourced soon

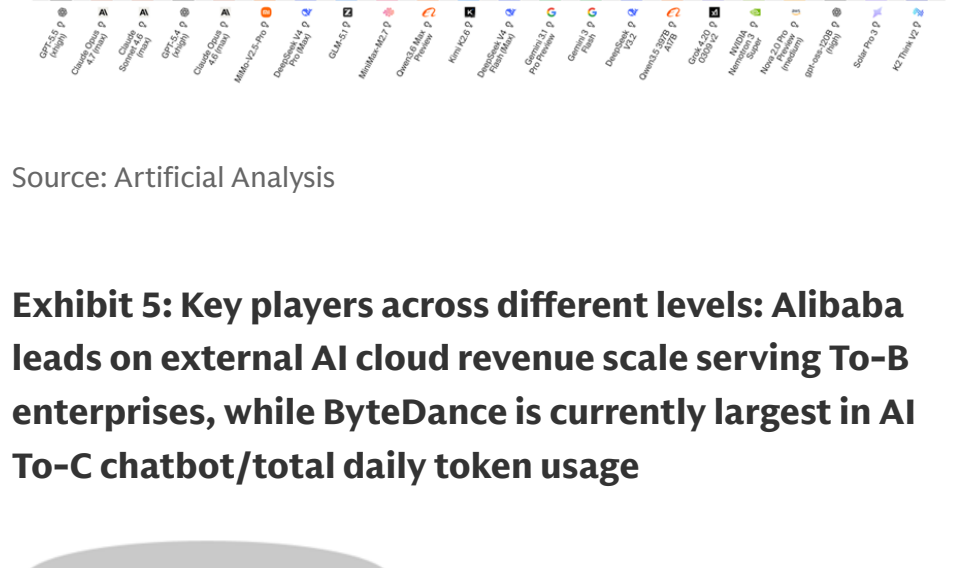

Source: Company data, OpenRouter, Data compiled by Goldman Sachs Global Investment Research

Exhibit 4: DeepSeek V4 Pro and V4 Flash leads amongst open-sourced models on GDPval-AA

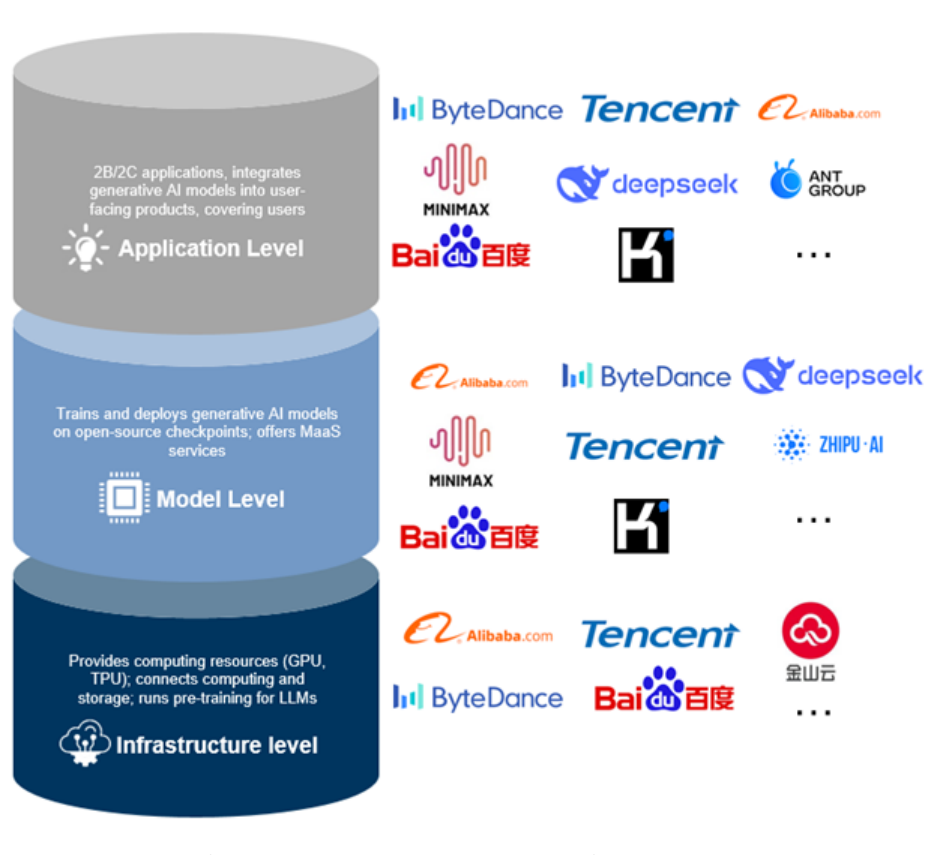

Source: Artificial Analysis

Exhibit 5: Key players across different levels: Alibaba leads on external AI cloud revenue scale serving To-B enterprises, while ByteDance is currently largest in AI To-C chatbot/total daily token usage

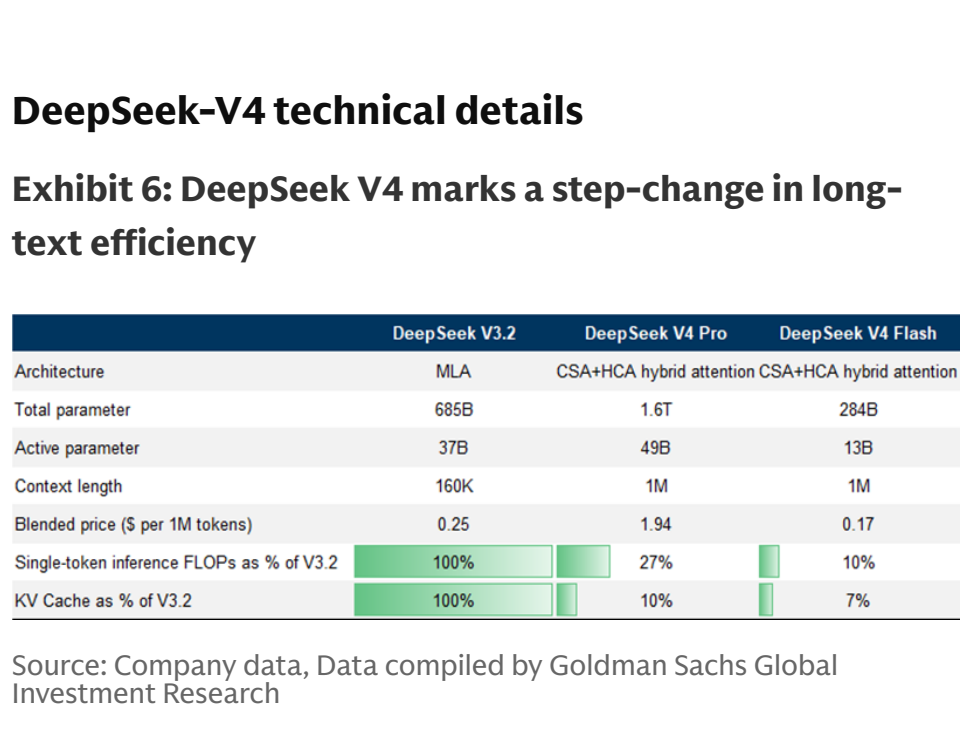
2B/2C applications, integrates generative AI models into user-facing products, covering users
Application Level
ByteDance
Tencent
Alibaba.com
MINIMAX
deepseek
ANT GROUP
Baidu 百度
…
Trains and deploys generative AI models on open-source checkpoints; offers MaaS services
Model Level
Alibaba.com
ByteDance
deepseek
MINIMAX
Tencent
ZHIPU·AI
Baidu 百度
…
Provides computing resources (GPU, TPU); connects computing and storage; runs pre-training for LLMs
Infrastructure level
Alibaba.com
Tencent
金山云
ByteDance
Baidu 百度
…

Source: Respective company data, Data compiled by Goldman Sachs Global Investment Research

DeepSeek-V4 technical details

Exhibit 6: DeepSeek V4 marks a step-change in long-text efficiency

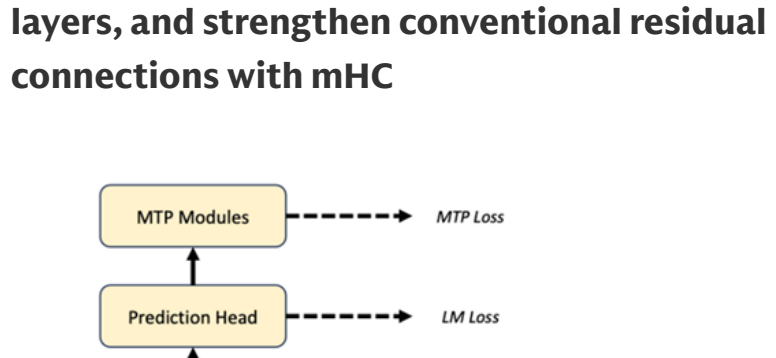

	DeepSeek V3.2	DeepSeek V4 Pro	DeepSeek V4 Flash
Architecture	MLA	CSA+HCA hybrid attention	CSA+HCA hybrid attention
Total parameter	685B	1.6T	284B
Active parameter	37B	49B	13B
Context length	160K	1M	1M
Blended price ($ per 1M tokens)	0.25	1.94	0.17
Single-token inference FLOPs as % of V3.2	100%	27%	10%
KV Cache as % of V3.2	100%	10%	7%

Source: Company data, Data compiled by Goldman Sachs Global Investment Research

Exhibit 7: Overall architecture of DeepSeek-V4 series, where DeepSeek used hybrid CSA (Compressed Sparse Attention) and HCA (Heavily Compressed Attention) for attention layers, DeepSeekMoE for feed-forward layers, and strengthen conventional residual connections with mHC

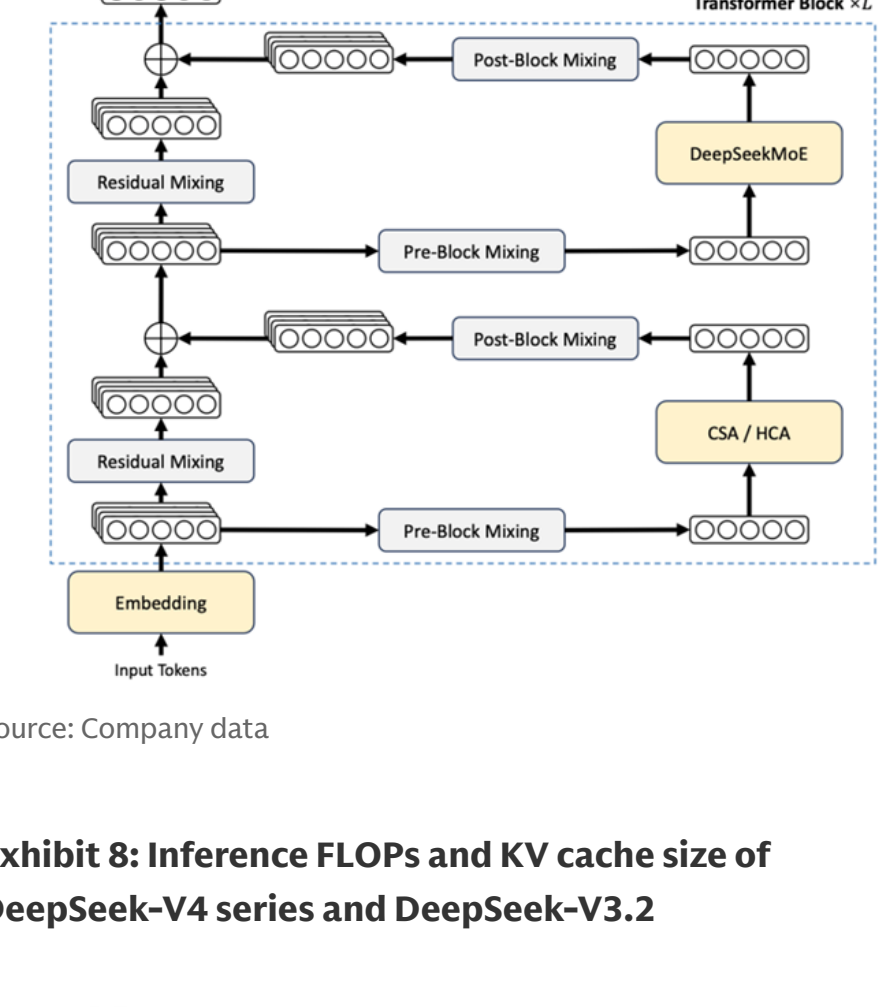

Source: Company data

Exhibit 8: Inference FLOPs and KV cache size of DeepSeek-V4 series and DeepSeek-V3.2

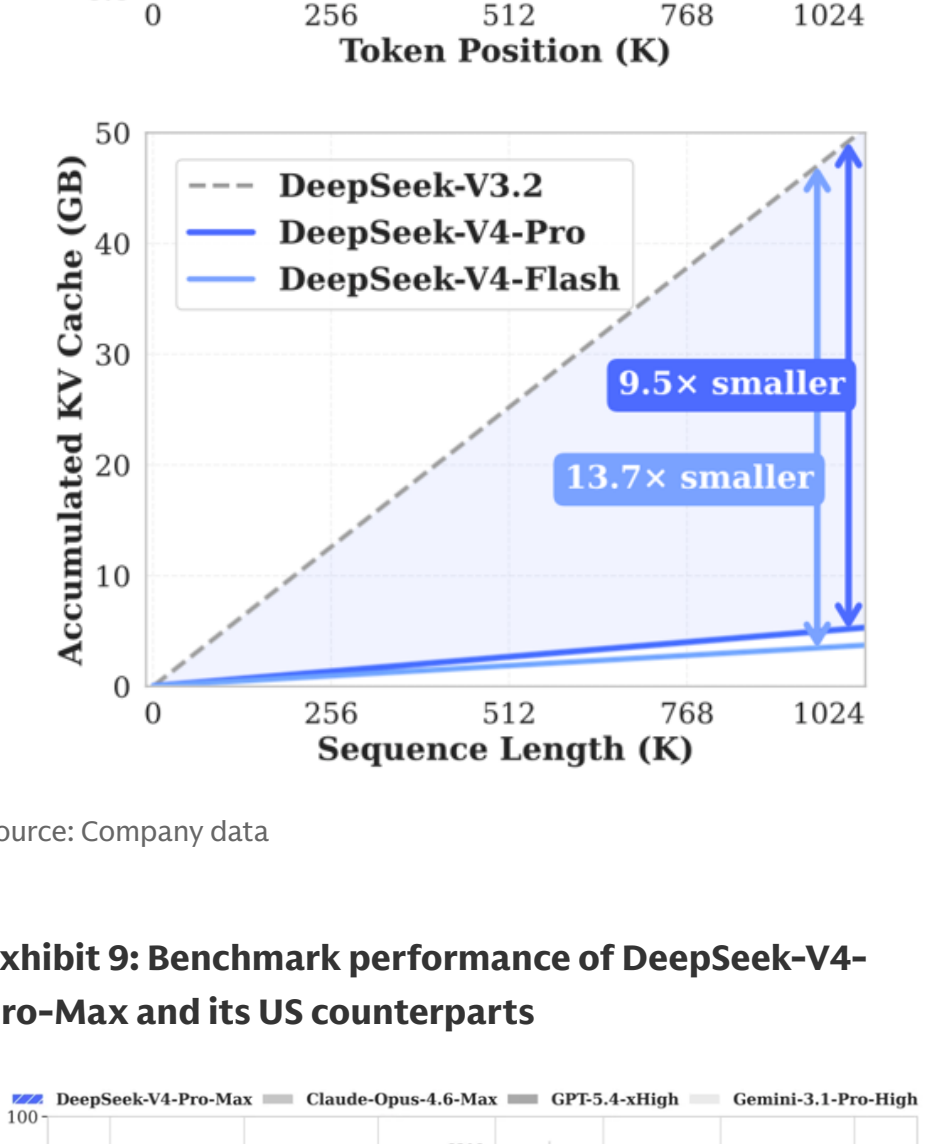

Source: Company data

Exhibit 9: Benchmark performance of DeepSeek-V4-Pro-Max and its US counterparts

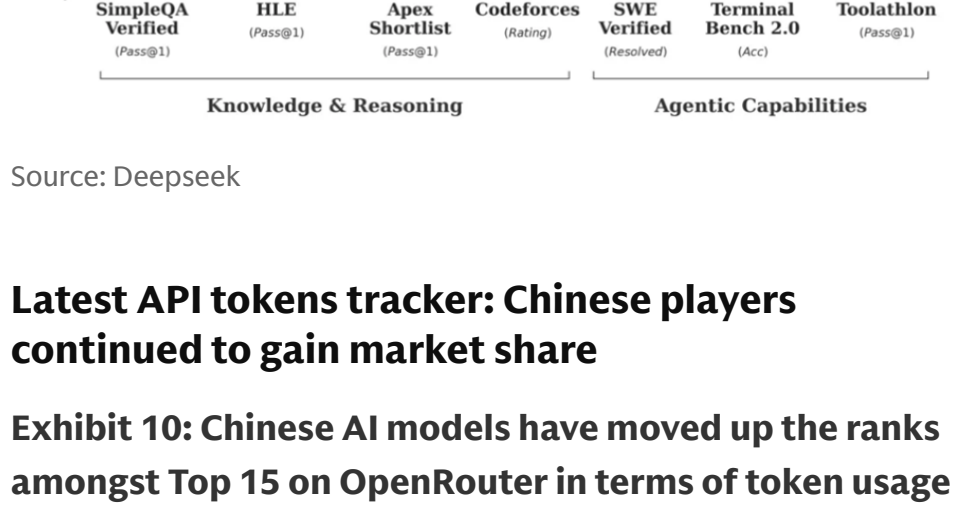

Source: Deepseek

Latest API tokens tracker: Chinese players continued to gain market share

Exhibit 10: Chinese AI models have moved up the ranks amongst Top 15 on OpenRouter in terms of token usage this month (Red denotes Chinese models)

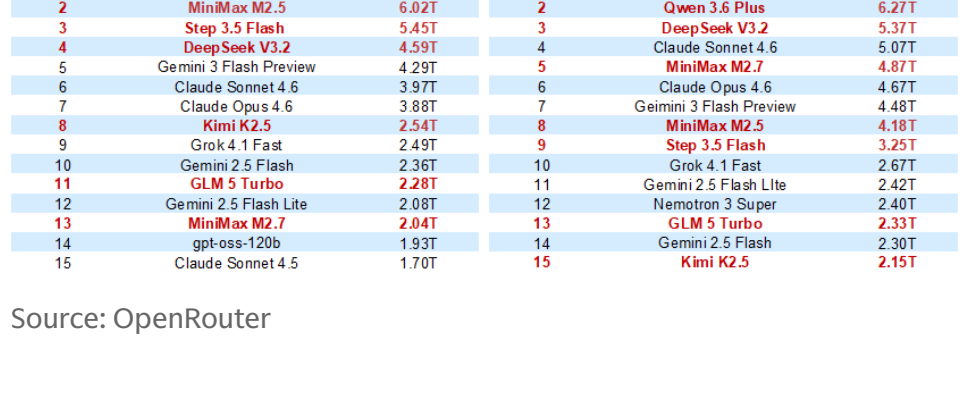

As of Apr 1st, 2026 (Rolling 30 days)

Ranking	Model	Tokens usage
1	MiMo-V2-Pro	6.87T
2	MiniMax M2.5	6.02T
3	Step 3.5 Flash	5.45T
4	DeepSeek V3.2	4.59T
5	Gemini 3 Flash Preview	4.29T
6	Claude Sonnet 4.6	3.97T
7	Claude Opus 4.6	3.88T
8	Kimi K2.5	2.54T
9	Grok 4.1 Fast	2.49T
10	Gemini 2.5 Flash	2.36T
11	GLM 5 Turbo	2.28T
12	Gemini 2.5 Flash Lite	2.08T
13	MiniMax M2.7	2.04T
14	gpt-oss-120b	1.93T
15	Claude Sonnet 4.5	1.70T

As of Apr 22nd, 2026 (Rolling 30 days)

Ranking	Model	Tokens usage
1	MiMo-V2-Pro	9.13T
2	Qwen 3.6 Plus	6.27T
3	DeepSeek V3.2	5.37T
4	Claude Sonnet 4.6	5.07T
5	MiniMax M2.7	4.87T
6	Claude Opus 4.6	4.67T
7	Geimini 3 Flash Preview	4.48T
8	MiniMax M2.5	4.18T
9	Step 3.5 Flash	3.25T
10	Grok 4.1 Fast	2.67T
11	Gemini 2.5 Flash LIte	2.42T
12	Nemotron 3 Super	2.40T
13	GLM 5 Turbo	2.33T
14	Gemini 2.5 Flash	2.30T
15	Kimi K2.5	2.15T

Source: OpenRouter

Exhibit 11: Chinese players have continued to gain market share through 2026

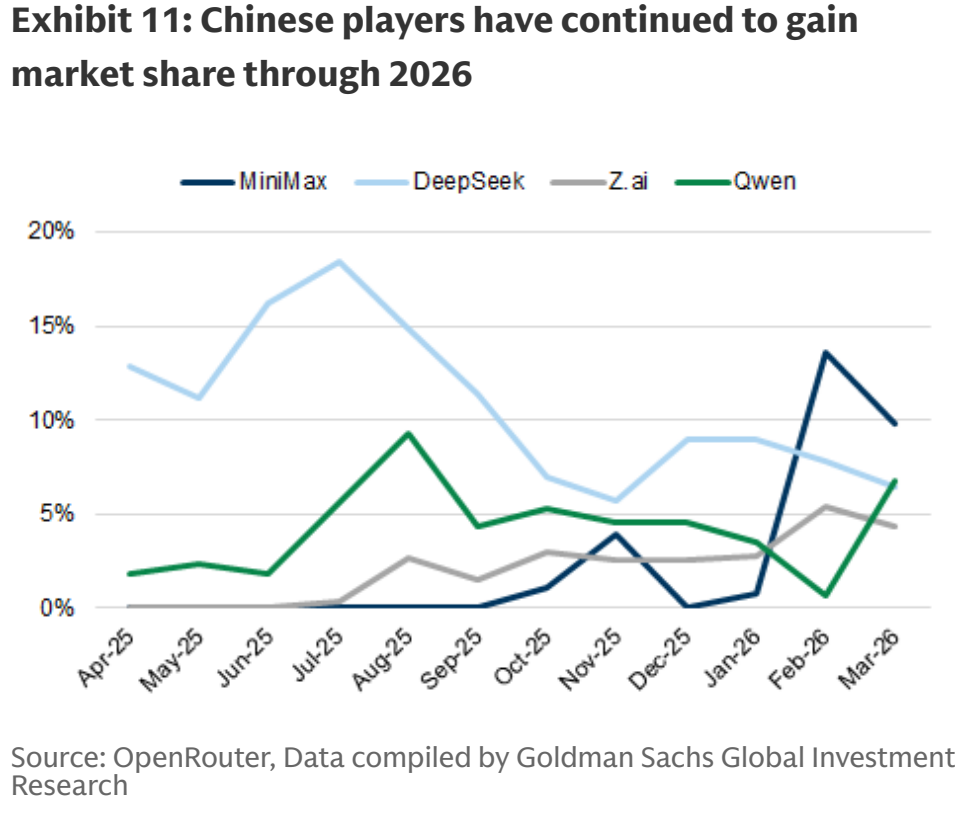

Source: OpenRouter, Data compiled by Goldman Sachs Global Investment Research